




Persistent

Selfie Shots - Tools 304

Kubernetes

Persistent University



Kubernetes Installation



Key learning points :

- Introduction
- Installing kubectl
- Exploring Kubernetes installation options
 - Google Kubernetes Engine
 - Minikube
 - kubeadm
 - hyperkube
 - Compiling from Source
- Installation of Pod Networks
- Deployment configurations

Introduction

This chapter is about Kubernetes installation and configuration. We are going to review a few installation mechanisms that you can use to create your own Kubernetes cluster. To get started without having to dive right away into installing and configuring a cluster, there are two main choices.

- One way is to use **Google Container Engine (GKE)**, a cloud service from the **Google Cloud Platform**, that lets you request a Kubernetes cluster with the latest stable version.
- Another easy way to get started is to use **Minikube**. It is a single binary which deploys into **Oracle VirtualBox** software, which can run in several operating systems. While Minikube is local and single node, it will give you a learning, testing, and development platform.
- In both cases, to be able to use the Kubernetes cluster, you will need to have installed the Kubernetes command line, called **kubectl**. This runs locally on your machine and targets the API server endpoint. It allows you to create, manage, and delete all Kubernetes resources (e.g. Pods, Deployments, Services). It is a powerful CLI that we will use throughout the rest of this course. So, you should become familiar with it.
- We will have a look at **kubeadm**, which is a newer tool coming up in the Kubernetes project, that makes installing Kubernetes easy and avoids vendor-specific installers. Getting a cluster running involves two commands: **kubeadm init**, that you run on a Master, and then, **kubeadm join**, that you run on your Worker Nodes, and your cluster bootstraps itself. The flexibility of these tools allows Kubernetes to be deployed in a number of places. We will be using this method on AWS nodes.
- We will also talk about other installation mechanisms that can be found in the community, such as **kubespray** or **kops**, another way to create a Kubernetes cluster on AWS. Additionally, you can use a container image called **hyperkube**, which contains all the key Kubernetes binaries, so that you can run a Kubernetes cluster by just starting a few containers on your nodes.

Installing kubectl

- To configure and manage your cluster, you will probably use the **kubectl** command. You can use **RESTful** calls or the **Go** language, as well.
- Enterprise Linux distributions have the various Kubernetes utilities and other files available in their repositories. For example, on RHEL 7/CentOS 7, you would find **kubectl** in the **kubernetes-client** package.
- You can (if needed) download the code from [Github](#), and go through the usual steps to compile and install **kubectl**.
- This command line will use **~/.kube/config** as a configuration file. This contains all the Kubernetes endpoints that you might use. If you examine it, you will see cluster definitions (i.e. IP endpoints), credentials, and contexts.
- A *context* is a combination of a cluster and user credentials. You can pass these parameters on the command line, or switch the shell between contexts with a command, as in:

\$ **kubectl config use-context foobar**

- This is handy when going from a local environment to a cluster in the cloud, or from one cluster to another, such as from development to production.

Using Google Kubernetes Engine (GKE)

- Google takes every Kubernetes release through rigorous testing and makes it available via its GKE service. To be able to use GKE, you will need the following:
- An account on Google Cloud.
- A method of payment for the services you will use.
- The gcloud command line client.
- There is an extensive documentation to get it installed. Pick your favorite method of installation and set it up. For more details, you can visit the [Installing Cloud SDK web page](#).
- You will then be able to follow the GKE quickstart guide and you will be ready to create your first Kubernetes cluster:

```
$ gcloud container clusters create test-cluster
```

```
$ gcloud container clusters list
```

```
$ kubectl get nodes
```

- By installing gcloud, you will have automatically installed kubectl. In the commands above, we created the cluster, listed it, and then, listed the nodes of the cluster with kubectl.
- Once you are done, do not forget to delete your cluster, otherwise you will keep on getting charged for it:

```
$ gcloud container clusters delete test-cluster
```

Using Minikube

- You can also use **Minikube**, an open source project within the **GitHub** [Kubernetes organization](https://github.com/kubernetes/minikube). While you can download a release from **GitHub**, following listed directions, it may be easier to download a pre-compiled binary. Make sure to verify and get the latest version.
- For example, to get the v0.22.2 version, do:

```
$ curl -Lo minikube https://storage.googleapis.com/minikube/releases/v0.22.2/minikube-linux-amd64
```

```
$ chmod +x minikube
```

```
$ sudo mv minikube /usr/local/bin
```

- With **Minikube** now installed, starting **Kubernetes** on your local machine is very easy:

```
$ minikube start
```

```
$ kubectl get nodes
```

- This will start a **VirtualBox** virtual machine that will contain a single node **Kubernetes** deployment and the **Docker** engine. Internally, **minikube** runs a single Go binary called **loalkube**. This binary runs all the components of Kubernetes together. This makes Minikube simpler than a full Kubernetes deployment. In addition, the Minikube VM also runs Docker, in order to be able to run containers.

Installing with kubeadm

- Once you become familiar with Kubernetes using Minikube, you may want to start building a real cluster. Currently, the most straightforward method is to use **kubeadm**, which appeared in Kubernetes v1.4.0, and can be used to bootstrap a cluster quickly.
- The Kubernetes website provides documentation on how to [use kubeadm to create a cluster](#).
- Package repositories are available for **Ubuntu 16.04** and **CentOS 7.1**.
- Run **kubeadm init** on the head node. The token is returned by the command, **kubeadm init**. Create a network for IP-per-Pod criteria. Run **kubeadm join --token token head-node-IP** on the worker nodes. You can also create the network with **kubectrl**, by using a resource manifest of the network.
- Once all the steps are completed, you will have a functional multi-node Kubernetes cluster, and you will be able to use **kubectrl** to interact with it.

Installing a Pod Network

- Prior to initializing the Kubernetes cluster, the network must be considered and IP conflicts avoided. There are several Pod networking choices, in varying levels of development and feature set:
- Calico
A flat Layer 3 network which communicates without IP encapsulation, used in production with software such as **Kubernetes**, **OpenShift**, **Docker**, **Mesos** and **OpenStack**. Viewed as a simple and flexible networking model, it scales well for large environments. Another network option, **Canal**, also part of this project, allows for integration with **Flannel**. Allows for implementation of network policies.
- Flannel
A Layer 3 IPv4 network between the nodes of a cluster. Developed by **CoreOS**, it has a long history with **Kubernetes**. Focused on traffic between hosts, not how containers configure local networking, it can use one of several backend mechanisms, such as VXLAN. A **flanneld** agent on each node allocates subnet leases for the host. While it can be configured after deployment, it is much easier prior to any Pods being added.
- Kube-router
Feature-filled single binary which claims to "*do it all*". The project is in the alpha stage, but promises to offer a distributed load balancer, firewall, and router purposely built for **Kubernetes**.
- Romana
Another project aimed at network and security automation for cloud native applications. Aimed at large clusters, IPAM-aware topology and integration with **kops** clusters.
- Weave Net
Typically used as an add-on for a CNI-enabled **Kubernetes** cluster.
- Many of the projects will mention the Container Network Interface (CNI), which is a CNCF project. Several container runtimes currently use CNI. As a standard to handle deployment management and cleanup of network resources, it will become more popular

More Installation Tools

- Since **Kubernetes** is, after all, like any other applications that you install on a server (whether physical or virtual), all the configuration management systems (e.g. **Chef**, **Puppet**, **Ansible**, **Terraform**) can be used. Various recipes are available on the Internet.
- Here are just a few examples of installation tools that you can use:
 - **kubespray** is now in the **Kubernetes** incubator. It is an advanced **Ansible** playbook which allows you to setup a **Kubernetes** cluster on various operating systems and use different network providers. It was once known as **kargo**.
 - **kops** lets you create a **Kubernetes** cluster on **AWS** via a single command line. Also in beta for **GKE** and alpha for **VMware**.
 - **kube-aws** is a command line tool that makes use of the **AWS Cloud Formation** to provision a **Kubernetes** cluster on **AWS**.
 - **kubicorn** is a tool which leverages the use of **kubeadm** to build a cluster. It claims to have no dependency on DNS, runs on several operating systems, and uses snapshots to capture a cluster and move it.

Installation Considerations

To begin the installation process, you should start experimenting with a single-node deployment. This single-node will run all the **Kubernetes** components (e.g. API server, controller, scheduler, kubelet, and kube-proxy). You can do this with **Minikube** for example.

Once you want to deploy on a cluster of servers (physical or virtual), you will have many choices to make, just like with any other distributed system:

- Which provider should I use? A public or private cloud? Physical or virtual?
- Which operating system should I use? **Kubernetes** runs on most operating systems (e.g. **Debian**, **Ubuntu**, **CentOS**, etc.), plus on container-optimized OSes (e.g. **CoreOS**, **Atomic**).
- Which networking solution should I use? Do I need an overlay?
- Where should I run my **etcd** cluster?
- Can I configure Highly Available (HA) head nodes?

To learn more about how to choose the best options, you can read the [Picking the Right Solution](#) article.

Deployment Configurations

- At a high level, you have four main deployment configurations:
 - Single-node
 - Single head node, multiple workers
 - Multiple head nodes with HA, multiple workers
 - HA **etcd**, HA head nodes, multiple workers.
- Which of the four you will use will depend on how advanced you are in your **Kubernetes** journey, but also on what your goals are.
- With a single-node deployment, all the components run on the same server. This is great for testing, learning, and developing around Kubernetes.
- Adding more workers, a single head node and multiple workers typically will consist of a single node **etcd** instance running on the head node with the API, the scheduler, and the controller-manager.
- Multiple head nodes in an HA configuration and multiple workers add more durability to the cluster. The API server will be fronted by a load balancer, the scheduler and the controller-manager will elect a leader (which is configured via flags). The **etcd** setup can still be single node.
- The most advanced and resilient setup would be an HA **etcd** cluster, with HA head nodes and multiple workers. Also, **etcd** would run as a true cluster, which would provide HA and would run on nodes separate from the Kubernetes head nodes.
- The use of Kubernetes Federations also offers high availability. Multiple clusters are joined together with a common control plane allowing movement of resources from one cluster to another administratively or after failure.

Using Hyperkube

- While you can run all the components as regular system daemons in unit files, you can also run the API server, the scheduler, and the controller-manager as containers. This is what **kubeadm** does.
- Indeed, there is a very handy all-in-one binary named **hyperkube**, which is available as a container image (e.g. **gcr.io/google_containers/hyperkube:v1.9.2**).
- This method of installation consists in running a **kubelet** as a system daemon and configuring it to read in manifests that specify how to run the other components (i.e. the API server, the scheduler, **etcd**, the controller). In these manifests, the **hyperkube** image is used. The **kubelet** will watch over them and make sure they get restarted if they die.
- To get a feel for this, you can simply download the **hyperkube** image and run a container to get help usage:

```
$ docker run --rm gcr.io/google_containers/hyperkube:v1.9.2 /hyperkube apiserver --help
```

```
$ docker run --rm gcr.io/google_containers/hyperkube:v1.9.2 /hyperkube scheduler --help
```

```
$ docker run --rm gcr.io/google_containers/hyperkube:v1.9.2 /hyperkube controller-manager --help
```

- This is also a very good way to start learning the various configuration flags.

Compiling from Source

- The [list of binary releases](#) is available on GitHub. Together with **gcloud**, **minikube**, and **kubeadmin**, these cover several scenarios to get started with **Kubernetes**.
- **Kubernetes** can also be compiled from source relatively quickly. You can clone the repository from **GitHub**, and then use the **Makefile** to build the binaries. You can build them natively on your platform if you have a Golang environment properly setup, or via **Docker** containers if you are on a [Docker host](#).
- To build natively with **Golang**, first install Golang. Download files and directions can be found online. <https://golang.org/doc/install>.
- Once Golang is working, you can clone the **kubernetes** repository, around 500MB in size. Change into the directory and use **make**:

```
$ cd $GOPATH
```

```
$ git clone https://github.com/kubernetes/kubernetes
```

```
$ cd kubernetes
```

```
$ make
```

- On a **Docker** host, clone the repository anywhere you want and use the **make quick-release** command. The build will be done in **Docker** containers.
- The **_output/bin** directory will contain the newly built binaries.

Summary

At the end of this session, we see that you are now able to

- Download installation and configuration tools.
- Install a Kubernetes master and grow a cluster.
- Configure a network solution for secure communications.
- Discuss highly-available deployment considerations.

Lab Exercise

- Install Kubernetes
- Deploy A Simple Application (e.g. nginx)